

Consistency Guard Report — Translate the following sentence to French, but do not use a...

Generated 2026-09-26T22:16:05.311Z

Evaluation ID: cmuiy80cw00mrq3hfp1s26eop

Executive summary

• Question (reasoning / adversarial): "Translate the following sentence to French, but do not use any French words: "Hello, how are you?""

- 1 model(s) evaluated with 2 repeated run(s) each on 9/26/2026, 10:15:52 PM.
- Average reliability: 35/100. Average consistency: 76/100.
- 0 conflict(s) detected (0 high severity).
- Results were produced in DEMO MODE with deterministic mock providers.

Methodology

- Each model answered the same prompt N times (controlled repeated runs).
- Responses were normalized and compared semantically (token Jaccard + claim overlap), never by exact string equality.
- Factual claims were extracted, cross-compared for conflicts, and checked against retrieved reference evidence.
- No reference answer was provided: accuracy figures are ESTIMATES from cross-model semantic agreement, not ground-truth accuracy.
- Reliability = $0.35 \cdot \text{accuracy} + 0.20 \cdot \text{consistency} + 0.15 \cdot \text{relevance} + 0.10 \cdot \text{completeness} + 0.10 \cdot \text{grounding} + 0.10 \cdot \text{instruction-following}$.

Reliability scores

- GPT-4o mini: reliability 35/100 (consistency 76, grounding 0, relevance 38, completeness 60, instruction-following 78, hallucination risk 2).

Conflict analysis

- No cross-model conflicts were detected.

Evidence

- Claim (unverified): "Quantitative estimates center around 30 units in most references." — evidence: NIST — Measurement and units reference (<https://www.nist.gov/measurement>)
- Claim (unverified): "The widely accepted position is that french plays a defining role, while translate modify the outcom" — evidence: Encyclopedia Britannica — General reference (<https://www.britannica.com/>)
- Claim (unverified): "A commonly cited definition states that following is characterized by consistent, repeatable behavio" — evidence: Stanford Encyclopedia of Philosophy — Definitions (<https://plato.stanford.edu/>)
- Claim (unverified): "This can vary depending on context, so treat specific figures as approximate." — evidence: W3C — Technical specifications (<https://www.w3.org/TR/>)
- Claim (unverified): "In summary, the key takeaway is the central claim stated above." — evidence: Python Software Foundation — Documentation (<https://docs.python.org/3/>)

Prompt quality

- Prompt quality score: 47/100.
- Issues: No clear role assigned to the model.; Limited context provided..

Model metrics

Model | Reliability | Consistency | Accuracy | Grounding | Relevance | Avg latency | Est. cost

GPT-4o mini | 35 | 76 | 0 (est.) | 0 | 38 | 776 ms | \$0.0001

Limitations

- Scores are estimates produced by heuristic analysis, not absolute measures of truth.
- Without a reference answer, "accuracy" reflects cross-model agreement, which can be wrong if all models err together.
- Evidence retrieval in this build uses a demo reference corpus; connect a live evidence source for production use.
- Cost figures are estimates from published per-token pricing and reported token usage.